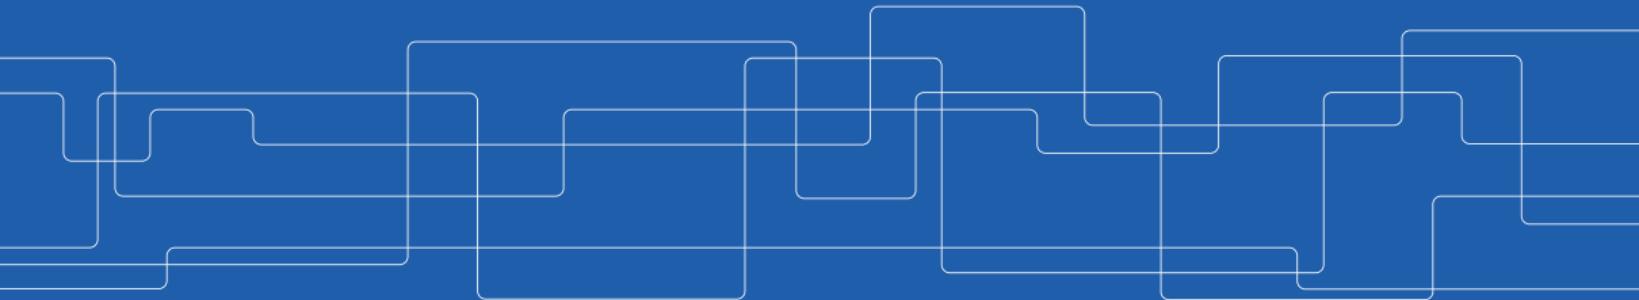




Processes - Part II

Amir H. Payberah
payberah@kth.se
2022





Threads

Thread


A basic unit of CPU utilization.



<https://tinyurl.com/e8crhtne>


Threads (1/2)

- A traditional process: has a single **thread**.




Threads (1/2)

- ▶ A traditional process: has a single **thread**.
- ▶ **Multiple threads** in a process: performing **more than one task** at a time.



Threads (1/2)

- ▶ A traditional process: has a single **thread**.
- ▶ **Multiple threads** in a process: performing **more than one task** at a time.
- ▶ Threads in a process **share** code section, data section, and **other OS resources**, e.g., open files.




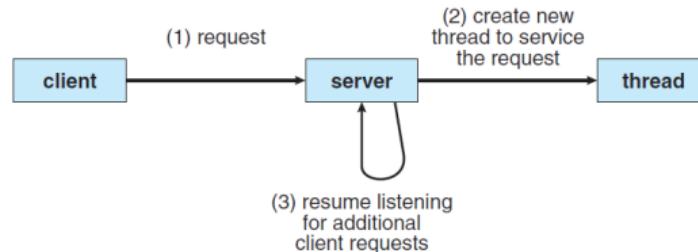
Threads (2/2)

- ▶ Multiple tasks of an application can be implemented by separate threads.
 - Update display
 - Fetch data
 - Spell checking
 - Answer a network request



Threads - Example

- ▶ Multi-threaded web-server architecture





Threads Benefits

- ▶ **Responsiveness:** allow **continued execution** if part of process is **blocked**.



Threads Benefits

- ▶ **Responsiveness:** allow **continued execution** if part of process is **blocked**.
- ▶ **Resource Sharing:** threads **share resources** of process, easier than **shared memory** or **message passing**.



Threads Benefits

- ▶ **Responsiveness:** allow **continued execution** if part of process is **blocked**.
- ▶ **Resource Sharing:** threads **share resources** of process, easier than **shared memory** or **message passing**.
- ▶ **Economy:** **thread switching** has lower overhead than **context switching**.



Threads Benefits

- ▶ **Responsiveness:** allow **continued execution** if part of process is **blocked**.
- ▶ **Resource Sharing:** threads **share resources** of process, easier than **shared memory** or **message passing**.
- ▶ **Economy:** **thread switching** has lower overhead than **context switching**.
- ▶ **Scalability:** process can take advantage of **multiprocessor architectures**.



Multi-core Programming




Multi-core Systems

- ▶ Users need **more computing performance**: single-CPU → multi-CPU

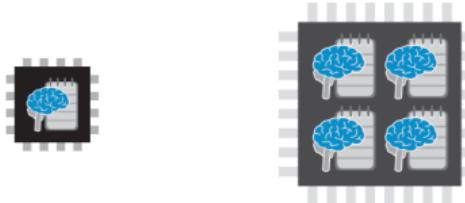
Multi-core Systems

- ▶ Users need **more computing performance**: single-CPU → multi-CPU
- ▶ A similar trend in system design: **multi-core** systems
 - Each **core** appears as a **separate processor**.



Multi-core Systems

- ▶ Users need **more computing performance**: single-CPU → multi-CPU
- ▶ A similar trend in system design: **multi-core** systems
 - Each **core** appears as a **separate processor**.



- ▶ **Multi-threaded programming**
 - Improves **concurrency** and more **efficient** use of multiple cores.



Concurrency vs. Parallelism (1/2)

- ▶ **Concurrency**: supporting **more than one task** by allowing all the **tasks to make progress**.




Concurrency vs. Parallelism (1/2)

- ▶ **Concurrency**: supporting **more than one task** by allowing all the **tasks to make progress**.
 - A **scheduler** providing concurrency.

Concurrency vs. Parallelism (1/2)

- ▶ **Concurrency**: supporting **more than one task** by allowing all the **tasks** to make progress.
 - A **scheduler** providing concurrency.
- ▶ **Concurrent execution** on a **single-core** system.



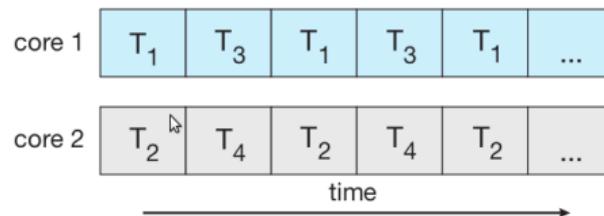


Concurrency vs. Parallelism (2/2)

- ▶ **Parallelism:** performing **more than one task simultaneously**.

Concurrency vs. Parallelism (2/2)


- ▶ **Parallelism:** performing **more than one task simultaneously**.
- ▶ **Parallelism on a multi-core system.**



Types of Parallelism

► Data parallelism


- Distributes subsets of the **same data** across multiple cores, **same operation** on each.



Types of Parallelism


► Data parallelism

- Distributes subsets of the **same data** across multiple cores, **same operation** on each.



► Task parallelism

- Distributes **threads** across cores, each thread performing **unique operation**.






Multi-threading Models

User Threads and Kernel Threads

- ▶ **User threads:** managed by [user-level threads library](#).




- ▶ **Kernel threads:** supported by the [Kernel](#).

User Threads and Kernel Threads

► **User threads:** managed by [user-level threads library](#).


- Three primary [thread libraries](#):
- POSIX pthreads
- Windows threads
- Java threads

► **Kernel threads:** supported by the [Kernel](#).



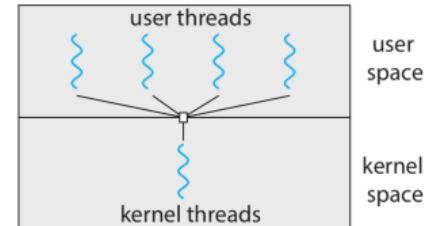
Multi-Threading Models

- ▶ Many-to-One
- ▶ One-to-One
- ▶ Many-to-Many



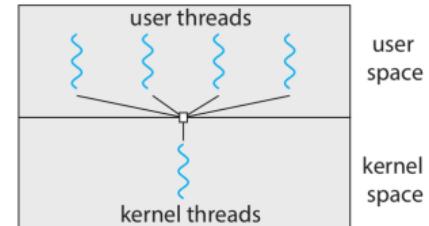
Many-to-One Model

- ▶ Many user-level threads mapped to single kernel thread.



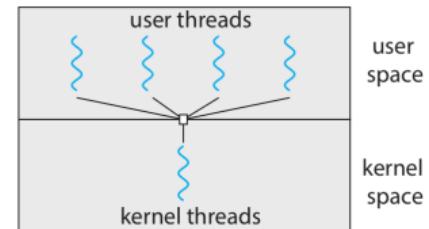
Many-to-One Model

- ▶ Many user-level threads mapped to single kernel thread.
- ▶ One thread blocking causes all to block.



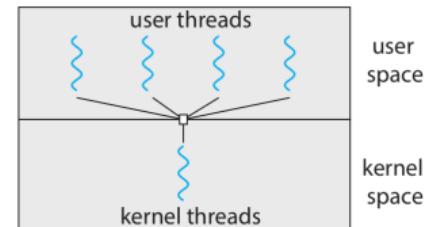
Many-to-One Model

- ▶ Many user-level threads mapped to single kernel thread.
- ▶ One thread blocking causes all to block.
- ▶ Multiple threads may not run in parallel on multi-core system because only one may be in kernel at a time.




Many-to-One Model

- ▶ Many user-level threads mapped to single kernel thread.
- ▶ One thread blocking causes all to block.
- ▶ Multiple threads may not run in parallel on multi-core system because only one may be in kernel at a time.
- ▶ Few systems currently use this model.
 - Solaris green threads
 - GNU portable threads



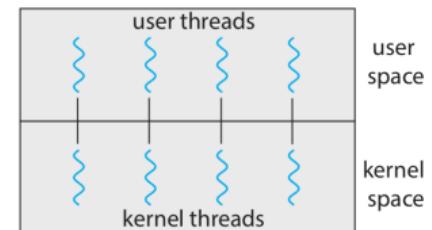
One-to-One Model

- ▶ Each user-level thread maps to one kernel thread.




One-to-One Model

- ▶ Each **user-level** thread maps to one **kernel** thread.
- ▶ Creating a **user-level** thread creates a **kernel** thread.



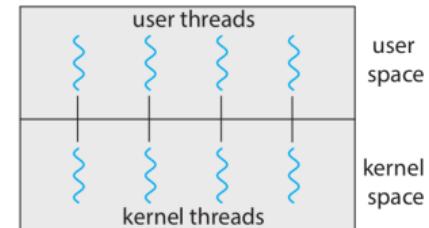
One-to-One Model

- ▶ Each **user-level** thread maps to **one kernel** thread.
- ▶ Creating a **user-level** thread creates a **kernel** thread.
- ▶ More **concurrency** than **many-to-one**.



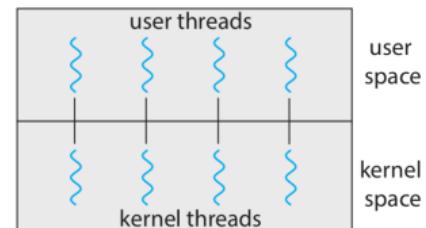
One-to-One Model

- ▶ Each user-level thread maps to one kernel thread.
- ▶ Creating a user-level thread creates a kernel thread.
- ▶ More concurrency than many-to-one.
- ▶ Number of threads per process sometimes restricted due to overhead.



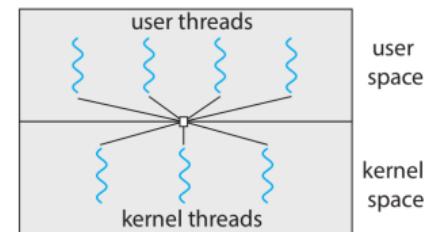
One-to-One Model

- ▶ Each user-level thread maps to one kernel thread.
- ▶ Creating a user-level thread creates a kernel thread.
- ▶ More concurrency than many-to-one.
- ▶ Number of threads per process sometimes restricted due to overhead.
- ▶ Examples:
 - Windows
 - Linux




Many-to-Many Model

- ▶ Allows **many user-level** threads to be mapped to **many kernel** threads.



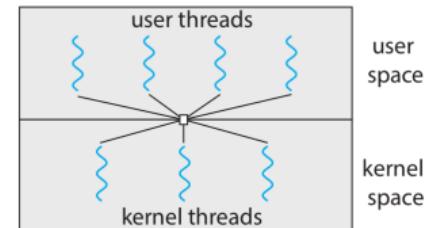
Many-to-Many Model

- ▶ Allows **many user-level** threads to be mapped to **many kernel** threads.
- ▶ Allows the OS to create a **sufficient number** of kernel threads.



Many-to-Many Model

- ▶ Allows **many user-level** threads to be mapped to **many kernel** threads.
- ▶ Allows the OS to create a **sufficient number** of kernel threads.
- ▶ Examples:
 - Windows with the ThreadFiber package
 - Otherwise not very common





Thread Libraries



Thread Libraries (1/2)

- ▶ Thread library provides programmer with API for creating and managing threads.



Thread Libraries (1/2)

- ▶ **Thread library** provides programmer with **API** for **creating and managing threads**.
- ▶ Two primary ways of implementing:
 - Library entirely in **user-space**.
 - **Kernel-level** library supported by the OS.



Thread Libraries (2/2)

► Pthread

- Either a user-level or a kernel-level library.



Thread Libraries (2/2)

- ▶ **Pthread**

- Either a user-level or a kernel-level library.

- ▶ **Windows thread**

- Kernel-level library.



Thread Libraries (2/2)

- ▶ **Pthread**


- Either a user-level or a kernel-level library.

- ▶ **Windows thread**

- Kernel-level library.

- ▶ **Java thread**

- Uses a thread library available on the host system.





Pthreads

- ▶ A POSIX API for **thread creation** and **synchronization**.



Pthreads

- ▶ A POSIX API for **thread creation** and **synchronization**.
- ▶ Specification, not implementation.



Pthreads

- ▶ A POSIX API for **thread creation** and **synchronization**.
- ▶ Specification, not implementation.
- ▶ API specifies **behavior** of the thread library, **implementation** is up to development of the library.



Pthreads

- ▶ A POSIX API for **thread creation** and **synchronization**.
- ▶ Specification, not implementation.
- ▶ API specifies **behavior** of the thread library, **implementation** is up to development of the library.
- ▶ Common in UNIX OSs, e.g., Solaris, Linux, Mac OS X



Thread ID

- ▶ The **thread ID (TID)** is the thread analogue to the process ID (PID).



Thread ID

- ▶ The **thread ID (TID)** is the thread analogue to the process ID (PID).
- ▶ The **PID** is assigned by the **Linux kernel**, and **TID** is assigned in the **Pthread library**.



Thread ID

- ▶ The **thread ID (TID)** is the thread analogue to the process ID (PID).
- ▶ The **PID** is assigned by the **Linux kernel**, and **TID** is assigned in the **Pthread library**.
- ▶ Represented by **pthread_t**.



Thread ID

- ▶ The **thread ID (TID)** is the thread analogue to the process ID (PID).
- ▶ The **PID** is assigned by the **Linux kernel**, and **TID** is assigned in the **Pthread library**.
- ▶ Represented by **pthread_t**.
- ▶ Obtaining a TID at runtime:

```
#include <pthread.h>  
  
pthread_t pthread_self(void);
```



Creating Threads

- ▶ `pthread_create()` defines and launches a new thread.

```
#include <pthread.h>

int pthread_create(pthread_t *thread, const pthread_attr_t *attr,
                  void *(*start_routine)(void *), void *arg);
```



Creating Threads

- ▶ `pthread_create()` defines and launches a new thread.

```
#include <pthread.h>

int pthread_create(pthread_t *thread, const pthread_attr_t *attr,
                  void *(*start_routine)(void *), void *arg);
```

- ▶ `start_routine` has the following signature:

```
void *start_thread(void *arg);
```



Terminating Threads

- ▶ Terminating yourself by calling `pthread_exit()`.

```
#include <pthread.h>

void pthread_exit(void *retval);
```



Terminating Threads

- ▶ Terminating yourself by calling `pthread_exit()`.

```
#include <pthread.h>

void pthread_exit(void *retval);
```

- ▶ Terminating others by calling `pthread_cancel()`.

```
#include <pthread.h>


int pthread_cancel(pthread_t thread);
```

Joining and Detaching Threads

- ▶ **Joining** allows one thread to **block** while **waiting for the termination** of another.

```
#include <pthread.h>

int pthread_join(pthread_t thread, void **retval);
int pthread_detach(pthread_t thread);
```




[<https://computing.llnl.gov/tutorials/pthreads/#Joining>]

Joining and Detaching Threads

- ▶ **Joining** allows one thread to **block** while **waiting for the termination** of another.
- ▶ You use **join** if you care about what value the thread returns when it is done, and use **detach** if you do not.

```
#include <pthread.h>

int pthread_join(pthread_t thread, void **retval);
int pthread_detach(pthread_t thread);
```



[<https://computing.llnl.gov/tutorials/pthreads/#Joining>]



A Threading Example

```
void *start_thread(void *message) {
    printf("%s\n", (const char *)message);
    return message;
}

int main(void) {
    pthread_t thread1, thread2;
    const char *message1 = "Thread 1";
    const char *message2 = "Thread 2";

    // Create two threads, each with a different message.
    pthread_create(&thread1, NULL, start_thread, (void *)message1);
    pthread_create(&thread2, NULL, start_thread, (void *)message2);

    // Wait for the threads to exit.
    pthread_join(thread1, NULL);
    pthread_join(thread2, NULL);

    return 0;
}
```



Implicit Threading



Implicit Threading

- ▶ Increasing the **number of threads**: program correctness more **difficult** with **explicit threads**.



Implicit Threading

- ▶ Increasing the **number of threads**: program correctness more **difficult** with **explicit threads**.
- ▶ **Implicit threading**: creation and management of threads done by **compilers and run-time libraries** rather than programmers.



Implicit Threading

- ▶ Increasing the **number of threads**: program correctness more **difficult** with **explicit threads**.
- ▶ **Implicit threading**: creation and management of threads done by **compilers and run-time libraries** rather than programmers.
- ▶ **Four** methods explored:
 - Thread Pools
 - Fork-Join
 - OpenMP
 - Grand Central Dispatch



Thread Pools

- ▶ Create a number of threads in a pool where they await work.



Thread Pools

- ▶ Create **a number of threads** in a pool where they await work.
- ▶ Usually slightly **faster** to service a request with an existing thread than **create a new thread**.

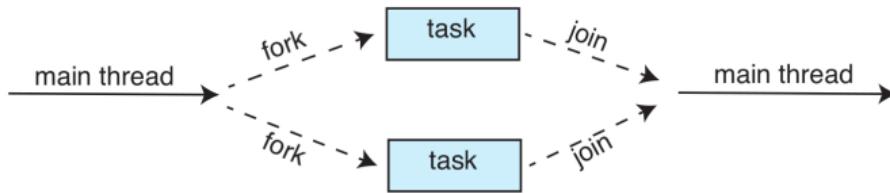


Thread Pools

- ▶ Create **a number of threads** in a pool where they **await work**.
- ▶ Usually slightly **faster** to service a request with an existing thread than **create a new thread**.
- ▶ Allows the number of threads in the application(s) to be **bound** to the size of the pool.

Fork-Join (1/2)

- ▶ Multiple threads (tasks) are forked, and then joined.





Fork-Join (2/2)

```
Task(problem)
    if problem is small enough
        solve the problem directly
    else
        subtask1 = fork(new Task(subset of problem))
        subtask2 = fork(new Task(subset of problem))

        result1 = join(subtask1)
        result2 = join(subtask2)


    return combined results
```

Fork-Join (2/2)

```
Task(problem)
    if problem is small enough
        solve the problem directly
    else
        subtask1 = fork(new Task(subset of problem))
        subtask2 = fork(new Task(subset of problem))

        result1 = join(subtask1)
        result2 = join(subtask2)

    return combined results
```





OpenMP (1/2)

- ▶ Set of compiler directives and APIs for C, C++, FORTRAN.



OpenMP (1/2)

- ▶ Set of compiler directives and APIs for C, C++, FORTRAN.
- ▶ Identifies parallel regions: blocks of code that can run in parallel.



OpenMP (1/2)

- ▶ Set of compiler directives and APIs for C, C++, FORTRAN.
- ▶ Identifies parallel regions: blocks of code that can run in parallel.
- ▶ `#pragma omp parallel`: create as many threads as there are cores.
- ▶ `#pragma omp parallel for`: run for loop in parallel.



OpenMP (2/2)

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[]) {

    /* sequential code */

    #pragma omp parallel
    {
        printf("I am a parallel region.");
    }

    /* sequential code */

    return 0;
}
```



Grand Central Dispatch

- ▶ Apple technology for Mac OS X and iOS: extensions to C, C++ API, and run-time library.



Grand Central Dispatch

- ▶ Apple technology for **Mac OS X and iOS**: extensions to C, C++ API, and run-time library.
- ▶ Allows identification of **parallel sections**.



Grand Central Dispatch

- ▶ Apple technology for Mac OS X and iOS: extensions to C, C++ API, and run-time library.
- ▶ Allows identification of parallel sections.
- ▶ Block is in ^{ }: ^{ printf("I am a block"); }



Grand Central Dispatch

- ▶ Apple technology for Mac OS X and iOS: extensions to C, C++ API, and run-time library.
- ▶ Allows identification of parallel sections.
- ▶ Block is in ^{ }: ^{ printf("I am a block"); }
- ▶ Blocks placed in dispatch queue.



Grand Central Dispatch

- ▶ Apple technology for Mac OS X and iOS: extensions to C, C++ API, and run-time library.
- ▶ Allows identification of parallel sections.
- ▶ Block is in ^{ }: ^{ printf("I am a block"); }
- ▶ Blocks placed in dispatch queue.

```
dispatch_queue_t queue = dispatch_get_global_queue(DISPATCH_QUEUE_PRIORITY_DEFAULT, 0);
dispatch_async(queue, ^{
    printf("I am a block.");
});
```



Threading Issues



Threading Issues

- ▶ The `fork()` and `exec()` system calls



Threading Issues

- ▶ The `fork()` and `exec()` system calls
- ▶ Signal handling



Threading Issues

- ▶ The `fork()` and `exec()` system calls
- ▶ Signal handling
- ▶ Thread cancellation



Threading Issues

- ▶ The `fork()` and `exec()` system calls
- ▶ Signal handling
- ▶ Thread cancellation
- ▶ Thread-Local Storage (TLS)



The fork() and exec() System Calls

- ▶ Does `fork()` duplicate only the **calling thread** or **all threads**?



The fork() and exec() System Calls

- ▶ Does `fork()` duplicate only the **calling thread** or **all threads**?
 - Some UNIXes have **two versions** of `fork`.



The fork() and exec() System Calls

- ▶ Does `fork()` duplicate only the **calling thread** or **all threads**?
 - Some UNIXes have **two versions** of `fork`.

- ▶ `exec()` usually works as **normal**: replace the **entire process**, including **all threads**.



Signal Handling (1/2)

- ▶ Signals are used in UNIX systems to **notify** a process that a particular event has occurred.



Signal Handling (1/2)

- ▶ Signals are used in UNIX systems to **notify** a process that a particular event has occurred.
- ▶ A **signal handler** is used to process signals:



Signal Handling (1/2)

- ▶ Signals are used in UNIX systems to **notify** a process that a **particular event has occurred**.
- ▶ A **signal handler** is used to process signals:
 1. Signal is **generated** by **particular event**.



Signal Handling (1/2)

- ▶ Signals are used in UNIX systems to notify a process that a particular event has occurred.
- ▶ A signal handler is used to process signals:
 1. Signal is generated by particular event.
 2. Signal is delivered to a process.



Signal Handling (1/2)

- ▶ Signals are used in UNIX systems to notify a process that a particular event has occurred.
- ▶ A signal handler is used to process signals:
 1. Signal is generated by particular event.
 2. Signal is delivered to a process.
 3. Signal is handled by the signal handlers, either the default or user-defined.



Signal Handling (1/2)

- ▶ Signals are used in UNIX systems to notify a process that a particular event has occurred.
- ▶ A signal handler is used to process signals:
 1. Signal is generated by particular event.
 2. Signal is delivered to a process.
 3. Signal is handled by the signal handlers, either the default or user-defined.
- ▶ Where should a signal be delivered for multi-threaded?



Signal Handling (2/2)

- ▶ Where should a signal be delivered for multi-threaded?



Signal Handling (2/2)

- ▶ Where should a signal be delivered for multi-threaded?
 - Deliver the signal to **the thread** to which the signal applies.



Signal Handling (2/2)

- ▶ Where should a signal be delivered for multi-threaded?
 - Deliver the signal to **the thread** to which the signal applies.
 - Deliver the signal to **every thread** in the process.



Signal Handling (2/2)

- ▶ Where should a signal be delivered for multi-threaded?
 - Deliver the signal to **the thread** to which the signal applies.
 - Deliver the signal to **every thread** in the process.
 - Deliver the signal to **certain threads** in the process.



Signal Handling (2/2)

- ▶ Where should a signal be delivered for multi-threaded?
 - Deliver the signal to **the thread** to which the signal applies.
 - Deliver the signal to **every thread** in the process.
 - Deliver the signal to **certain threads** in the process.
 - Assign a **specific thread** to receive **all signals** for the process.



Thread Cancellation

- ▶ Terminating a thread **before it has finished**.



Thread Cancellation

- ▶ Terminating a thread **before it has finished**.
- ▶ Thread to be **canceled** is **target thread**.



Thread Cancellation

- ▶ Terminating a thread **before it has finished**.
- ▶ Thread to be **canceled** is **target thread**.
- ▶ Two general approaches:



Thread Cancellation

- ▶ Terminating a thread **before it has finished**.
- ▶ Thread to be **canceled** is **target thread**.
- ▶ Two general approaches:
 - **Asynchronous cancellation** terminates the **target thread immediately**.



Thread Cancellation

- ▶ Terminating a thread **before it has finished**.
- ▶ Thread to be **canceled** is **target thread**.
- ▶ Two general approaches:
 - **Asynchronous cancellation** terminates the **target thread** immediately.
 - **Deferred cancellation** allows the **target thread** to **periodically check** if it should be cancelled.



Thread-Local Storage (TLS)

- ▶ TLS allows **each thread** to have **its own copy of data**.



Thread-Local Storage (TLS)

- ▶ TLS allows each thread to have its own copy of data.
- ▶ Useful when you do not have control over the thread creation process (i.e., thread pool)



Thread-Local Storage (TLS)

- ▶ TLS allows each thread to have its own copy of data.
- ▶ Useful when you do not have control over the thread creation process (i.e., thread pool)
- ▶ Different from local variables:



Thread-Local Storage (TLS)

- ▶ TLS allows each thread to have its own copy of data.
- ▶ Useful when you do not have control over the thread creation process (i.e., thread pool)
- ▶ Different from local variables:
 - Local variables visible only during single function invocation.



Thread-Local Storage (TLS)

- ▶ TLS allows each thread to have its own copy of data.
- ▶ Useful when you do not have control over the thread creation process (i.e., thread pool)
- ▶ Different from local variables:
 - Local variables visible only during single function invocation.
 - TLS visible across function invocations.



Summary



Summary

- ▶ Single-thread vs. Multi-thread



Summary

- ▶ Single-thread vs. Multi-thread
- ▶ Concurrency vs. parallelism



Summary

- ▶ Single-thread vs. Multi-thread
- ▶ Concurrency vs. parallelism
- ▶ Multi-threading models: many-to-one, one-to-one, many-to-many



Summary

- ▶ Single-thread vs. Multi-thread
- ▶ Concurrency vs. parallelism
- ▶ Multi-threading models: many-to-one, one-to-one, many-to-many
- ▶ Multi-thread libraries: pthread



Summary

- ▶ Single-thread vs. Multi-thread
- ▶ Concurrency vs. parallelism
- ▶ Multi-threading models: many-to-one, one-to-one, many-to-many
- ▶ Multi-thread libraries: pthread
- ▶ Implicit threading



Summary

- ▶ Single-thread vs. Multi-thread
- ▶ Concurrency vs. parallelism
- ▶ Multi-threading models: many-to-one, one-to-one, many-to-many
- ▶ Multi-thread libraries: pthread
- ▶ Implicit threading
- ▶ Threading issues



Questions?

Acknowledgements

Some slides were derived from Avi Silberschatz slides.