



Processes Synchronization - Part II

Amir H. Payberah
payberah@kth.se
2022



Deadlocks



Motivation

- ▶ Multiprogramming environment: several processes compete for a finite number of resources.



Motivation

- ▶ Multiprogramming environment: several processes compete for a finite number of resources.
- ▶ A process requests resources: if the resources are not available at that time, the process enters a waiting state.



Motivation

- ▶ **Multiprogramming** environment: **several processes** compete for a **finite number of resources**.
- ▶ A process requests resources: if the resources are **not available** at that time, the process enters a **waiting state**.
- ▶ What if the requests resources are **held by other waiting processes**?



Motivation

- ▶ **Multiprogramming** environment: **several processes** compete for a **finite number of resources**.
- ▶ A process requests resources: if the resources are **not available** at that time, the process enters a **waiting state**.
- ▶ What if the requests resources are **held by other waiting processes**?
- ▶ This situation is called a **deadlock**.



Deadlock System Model

- ▶ System consists of m resources: R_1, R_2, \dots, R_m



Deadlock System Model

- ▶ System consists of m resources: R_1, R_2, \dots, R_m
- ▶ Resource types: CPU cycles, memory space, I/O devices



Deadlock System Model

- ▶ System consists of m resources: R_1, R_2, \dots, R_m
- ▶ Resource types: CPU cycles, memory space, I/O devices
- ▶ Each resource type R_i has W_i instances.



Deadlock System Model

- ▶ System consists of m resources: R_1, R_2, \dots, R_m
- ▶ Resource types: CPU cycles, memory space, I/O devices
- ▶ Each resource type R_i has W_i instances.
- ▶ Each process utilizes a resource as follows:
 - Request
 - Use
 - Release



Deadlock Characterization (1/3)

- ▶ Deadlock can arise if **four conditions** hold **simultaneously**:
 - Mutual exclusion
 - Hold and wait
 - No preemption
 - Circular wait



Deadlock Characterization (2/3)

► Mutual exclusion

- Only one process at a time can use a resource.



Deadlock Characterization (2/3)

► Mutual exclusion

- Only one process at a time can use a resource.

► Hold and wait

- A process holding at least one resource is waiting to acquire additional resources held by other processes.



Deadlock Characterization (3/3)

► No preemption

- A resource can be released only **voluntarily** by the process holding it, after that process has **completed its task**.

Deadlock Characterization (3/3)

► No preemption

- A resource can be released only **voluntarily** by the process holding it, after that process has **completed its task**.

► Circular wait

- A set processes: $\{P_0, P_1, \dots, P_n\}$
- P_0 is **waiting** for a resource that is held by P_1
- P_1 is **waiting** for a resource that is held by P_2
- ...
- P_n is **waiting** for a resource that is held by P_0



Deadlock Example (1/2)

```
/* Create and initialize the mutex locks */  
pthread_mutex_t first_mutex;  
pthread_mutex_t second_mutex;  
  
pthread_mutex_init(&first_mutex, NULL);  
pthread_mutex_init(&second_mutex, NULL);
```




Deadlock Example (2/2)

```
void *thread_one(void *args) {  
    pthread_mutex_lock(&first_mutex);  
    pthread_mutex_lock(&second_mutex);  
    // do some work  
    pthread_mutex_unlock(&second_mutex);  
    pthread_mutex_unlock(&first_mutex);  
  
    pthread_exit(0);  
}
```

Deadlock Example (2/2)

```
void *thread_one(void *args) {  
    pthread_mutex_lock(&first_mutex);  
    pthread_mutex_lock(&second_mutex);  
    // do some work  
    pthread_mutex_unlock(&second_mutex);  
    pthread_mutex_unlock(&first_mutex);  
  
    pthread_exit(0);  
}
```

```
void *thread_two(void *args) {  
    pthread_mutex_lock(&second_mutex);  
    pthread_mutex_lock(&first_mutex);  
    // do some work  
    pthread_mutex_unlock(&first_mutex);  
    pthread_mutex_unlock(&second_mutex);  
  
    pthread_exit(0);  
}
```

Resource-Allocation Graph



Resource-Allocation Graph (1/2)

- ▶ A set of vertices V and a set of edges E .



Resource-Allocation Graph (1/2)

- ▶ A set of **vertices** V and a set of **edges** E .
- ▶ **Vertices**
 - All the **processes** in the system: $P = P_1, P_2, \dots, P_n$
 - All **resource types** in the system: $R = R_1, R_2, \dots, R_m$

Resource-Allocation Graph (1/2)

- ▶ A set of **vertices** V and a set of **edges** E .

- ▶ **Vertices**
 - All the **processes** in the system: $P = P_1, P_2, \dots, P_n$
 - All **resource types** in the system: $R = R_1, R_2, \dots, R_m$

- ▶ **Edges**
 - **Request edge**: directed edge $P_i \rightarrow R_j$
 - **Assignment edge**: directed edge $R_j \rightarrow P_i$

Resource-Allocation Graph (2/2)

► Process (vertices)



Resource-Allocation Graph (2/2)

► Process (vertices)



► Resource type with 4 instances (vertices)



Resource-Allocation Graph (2/2)

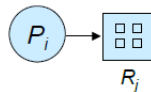
- ▶ Process (vertices)



- ▶ Resource type with 4 instances (vertices)



- ▶ P_i requests instance of R_j (edge)



-

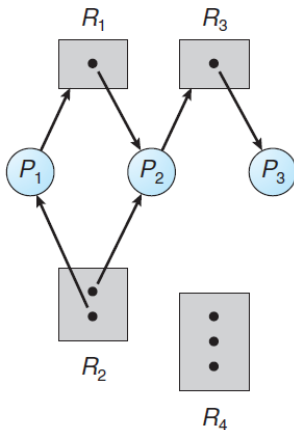
- 

-

-
- A diagram showing a node P_i (represented by a circle) connected to a set of nodes R_j (represented by a square containing a 2x2 grid of smaller squares). An arrow points from one of the squares in the R_j grid to the node P_i .


Resource-Allocation Graph Example (1/3)

- Example of a **resource allocation graph**.



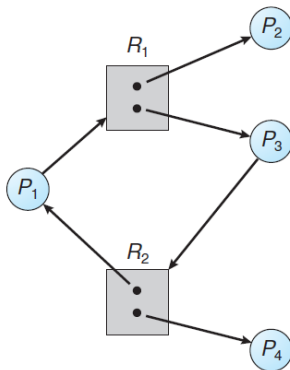
Resource-Allocation Graph Example (2/3)

- Resource allocation graph **with a deadlock**.



Resource-Allocation Graph Example (3/3)

- Resource allocation graph with a **cycle** but no deadlock.





Basic Facts

- ▶ If graph contains no cycles
 - No deadlock

Basic Facts

- ▶ If graph contains no cycles
 - No deadlock

- ▶ If graph contains a cycle
 - If only one instance per resource type, then deadlock.
 - If several instances per resource type, possibility of deadlock.



Methods for Handling Deadlocks

- ▶ Ensure that the system will **never enter a deadlock state**:



Methods for Handling Deadlocks

- ▶ Ensure that the system will **never enter a deadlock state**:
 - **Deadlock prevention**
 - **Deadlock avoidance**



Methods for Handling Deadlocks

- ▶ Ensure that the system will **never enter a deadlock state**:
 - **Deadlock prevention**
 - **Deadlock avoidance**

- ▶ **Allow** the system to **enter a deadlock** state and then **recover**.

Methods for Handling Deadlocks

- ▶ Ensure that the system will **never enter a deadlock state**:
 - **Deadlock prevention**
 - **Deadlock avoidance**

- ▶ **Allow** the system to **enter a deadlock** state and then **recover**.

- ▶ **Ignore the problem** and pretend that deadlocks never occur in the system; used by **most operating systems**.

Deadlock Prevention



Deadlock Prevention (1/3)

- ▶ Deadlock can arise if **four conditions** hold **simultaneously**:
 - Mutual exclusion
 - Hold and wait
 - No preemption
 - Circular wait



Deadlock Prevention (1/3)

- ▶ Deadlock can arise if **four conditions** hold **simultaneously**:
 - Mutual exclusion
 - Hold and wait
 - No preemption
 - Circular wait

- ▶ **Restrain** the ways requests can be made.



Deadlock Prevention (2/3)

► Mutual exclusion

- Not required for **sharable** resources, e.g., **read-only files**.
- Must hold for **non-sharable** resources.



Deadlock Prevention (2/3)

► Mutual exclusion

- Not required for **sharable** resources, e.g., **read-only files**.
- Must hold for **non-sharable** resources.

► Hold and wait

- Must guarantee that whenever a process requests a resource, it **does not hold any other processes**.

Deadlock Prevention (2/3)

► Mutual exclusion

- Not required for **sharable** resources, e.g., **read-only files**.
- Must hold for **non-sharable** resources.

► Hold and wait

- Must guarantee that whenever a process requests a resource, it **does not hold any other processes**.
- **Solution 1**: require a process to request and be **allocated all its resources before it begins execution**.

Deadlock Prevention (2/3)

► Mutual exclusion

- Not required for **sharable** resources, e.g., **read-only files**.
- Must hold for **non-sharable** resources.

► Hold and wait

- Must guarantee that whenever a process requests a resource, it **does not hold any other processes**.
- **Solution 1**: require a process to request and be **allocated all its resources before it begins execution**.
- **Solution 2**: allows a process to request resources only when **it has none**.

Deadlock Prevention (2/3)

► Mutual exclusion

- Not required for **sharable** resources, e.g., **read-only files**.
- Must hold for **non-sharable** resources.

► Hold and wait

- Must guarantee that whenever a process requests a resource, it **does not hold any other processes**.
- **Solution 1**: require a process to request and be **allocated all its resources before it begins execution**.
- **Solution 2**: allows a process to request resources only when **it has none**.
- **Low resource utilization**
- **Starvation** possible



Deadlock Prevention (3/3)

- ▶ No preemption



Deadlock Prevention (3/3)

► No preemption

- If a process that is holding some resources, requests another resource that **cannot be immediately allocated** to it, then all resources currently being held are **released**.



Deadlock Prevention (3/3)

► No preemption

- If a process that is holding some resources, requests another resource that **cannot be immediately allocated** to it, then all resources currently being held are **released**.
- **Preempted resources** are added to the list of resources for which the process is waiting.

Deadlock Prevention (3/3)

► No preemption

- If a process that is holding some resources, requests another resource that **cannot be immediately allocated** to it, then all resources currently being held are **released**.
- **Preempted resources** are added to the list of resources for which the process is waiting.
- Process will be restarted only when it can **regain its old resources**, as well as the **new ones** that it is requesting.

Deadlock Prevention (3/3)

► No preemption

- If a process that is holding some resources, requests another resource that **cannot be immediately allocated** to it, then all resources currently being held are **released**.
- **Preempted resources** are added to the list of resources for which the process is waiting.
- Process will be restarted only when it can **regain its old resources**, as well as the **new ones** that it is requesting.


► Circular wait

- Impose a **total ordering** of all resource types, and require that each process **requests resources in an increasing order** of enumeration.

Deadlock Avoidance


Basic Facts

- ▶ If a system is in the **safe state**
 - **No deadlock**




Basic Facts

- ▶ If a system is in the **safe state**
 - **No deadlock**
- ▶ If a system is in the **unsafe state**
 - **Possibility of deadlock**



Basic Facts

- ▶ If a system is in the **safe state**
 - **No deadlock**
- ▶ If a system is in the **unsafe state**
 - **Possibility of deadlock**
- ▶ **Avoidance**
 - Ensure that a system will never enter an unsafe state.





Safe State (1/2)

- ▶ When a process requests an **available resource**, system must decide if immediate allocation leaves the system in a **safe state**.

Safe State (1/2)

- ▶ When a process requests an **available resource**, system must decide if immediate allocation leaves the system in a **safe state**.
- ▶ **Safe state**: there exists a sequence $\langle P_1, P_2, \dots, P_n \rangle$ of **all** the processes in the systems such that for each P_i , the resources that P_i **can still request** be satisfied by:
currently available resources + resources held by all the P_j , with $j < i$.

Safe State (2/2)

- ▶ If P_i resource needs are **not immediately available**, then P_i can **wait** until all P_j have finished.

Safe State (2/2)

- ▶ If P_i resource needs are **not immediately available**, then P_i can **wait** until all P_j have finished.
- ▶ When P_j is **finished**, P_i can **obtain needed resources**, execute, return allocated resources, and terminate.

Safe State (2/2)

- ▶ If P_i resource needs are **not immediately available**, then P_i can **wait** until all P_j have finished.
- ▶ When P_j is **finished**, P_i can **obtain needed resources**, execute, return allocated resources, and terminate.
- ▶ When P_i **terminates**, P_{i+1} can obtain its needed resources, and so on.

Safe Mode Example (1/2)

- ▶ 3 processes: P_0 through P_2
- ▶ 1 resource type:
 - A (12 instances)
- ▶ Snapshot at time T_0

	<u>Maximum Needs</u>	<u>Current Needs</u>
P_0	10	5
P_1	4	2
P_2	9	2

Safe Mode Example (1/2)

- ▶ 3 processes: P_0 through P_2
- ▶ 1 resource type:
 - A (12 instances)
- ▶ Snapshot at time T_0

	<u>Maximum Needs</u>	<u>Current Needs</u>
P_0	10	5
P_1	4	2
P_2	9	2

- ▶ Safe mode sequence?

Safe Mode Example (1/2)

- ▶ 3 processes: P_0 through P_2
- ▶ 1 resource type:
 - A (12 instances)
- ▶ Snapshot at time T_0

	<u>Maximum Needs</u>	<u>Current Needs</u>
P_0	10	5
P_1	4	2
P_2	9	2

- ▶ Safe mode sequence? $\langle P_1, P_0, P_2 \rangle$

Safe Mode Example (2/2)

- ▶ 3 processes: P_0 through P_2
- ▶ 1 resource type:
 - A (12 instances)
- ▶ Snapshot at time T_0

	<u>Maximum Needs</u>	<u>Current Needs</u>
P_0	10	5
P_1	4	2
P_2	9	2

Safe Mode Example (2/2)

- ▶ 3 processes: P_0 through P_2
- ▶ 1 resource type:
 - A (12 instances)
- ▶ Snapshot at time T_0

	<u>Maximum Needs</u>	<u>Current Needs</u>
P_0	10	5
P_1	4	2
P_2	9	2

- ▶ Suppose that, at time T_1 , process P_2 requests and is allocated one more resource.

Safe Mode Example (2/2)

- ▶ 3 processes: P_0 through P_2
- ▶ 1 resource type:
 - A (12 instances)
- ▶ Snapshot at time T_0

	<u>Maximum Needs</u>	<u>Current Needs</u>
P_0	10	5
P_1	4	2
P_2	9	2

- ▶ Suppose that, at time T_1 , process P_2 requests and is allocated one more resource.
- ▶ Safe mode sequence?

Safe Mode Example (2/2)

- ▶ 3 processes: P_0 through P_2
- ▶ 1 resource type:
 - A (12 instances)
- ▶ Snapshot at time T_0

	<u>Maximum Needs</u>	<u>Current Needs</u>
P_0	10	5
P_1	4	2
P_2	9	2

- ▶ Suppose that, at time T_1 , process P_2 requests and is allocated one more resource.
- ▶ Safe mode sequence? Not safe



Avoidance Algorithms

- ▶ Single instance of a resource type
 - Use a resource-allocation graph



Avoidance Algorithms

- ▶ Single instance of a resource type
 - Use a resource-allocation graph
- ▶ Multiple instances of a resource type
 - Use the banker's algorithm

Resource-Allocation Graph Algorithm



Resource-Allocation Graph Scheme

- **Claim edge** $P_i \rightarrow R_j$: indicates that process P_j may **request** resource R_j ; represented by a dashed line



Resource-Allocation Graph Scheme

- ▶ **Claim edge** $P_i \rightarrow R_j$: indicates that process P_j may **request** resource R_j ; represented by a dashed line
- ▶ **Claim edge** converts to **request edge** when a process **requests a resource**.

Resource-Allocation Graph Scheme

- ▶ **Claim edge** $P_i \rightarrow R_j$: indicates that process P_j may **request** resource R_j ; represented by a dashed line
- ▶ **Claim edge** converts to **request edge** when a process **requests a resource**.
- ▶ **Request edge** converted to an **assignment edge** when the **resource is allocated** to the process.


Resource-Allocation Graph Scheme

- ▶ **Claim edge** $P_i \rightarrow R_j$: indicates that process P_j may **request** resource R_j ; represented by a dashed line
- ▶ **Claim edge** converts to **request edge** when a process **requests a resource**.
- ▶ **Request edge** converted to an **assignment edge** when the **resource is allocated** to the process.
- ▶ When a resource is **released** by a process, **assignment edge** reconverts to a **claim edge**.


Resource-Allocation Graph Scheme

- ▶ **Claim edge** $P_i \rightarrow R_j$: indicates that process P_j may **request** resource R_j ; represented by a dashed line
- ▶ **Claim edge** converts to **request edge** when a process **requests a resource**.
- ▶ **Request edge** converted to an **assignment edge** when the **resource is allocated** to the process.
- ▶ When a resource is **released** by a process, **assignment edge** reconverts to a **claim edge**.
- ▶ Resources must be claimed **a priori** in the system.

Resource-Allocation Graph



Unsafe State In Resource-Allocation Graph





Resource-Allocation Graph Algorithm

- ▶ Suppose that process P_i requests a resource R_j .
- ▶ The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph.

Banker's Algorithm



Banker's Algorithm

- ▶ Multiple instances



Banker's Algorithm

- ▶ Multiple instances
- ▶ Each process must **a priori claim** of the **maximum** use.



Banker's Algorithm

- ▶ Multiple instances
- ▶ Each process must **a priori claim** of the **maximum** use.
- ▶ When a process **requests a resource** it may have to **wait**.



Banker's Algorithm

- ▶ Multiple instances
- ▶ Each process must **a priori claim** of the **maximum** use.
- ▶ When a process **requests a resource** it may have to **wait**.
- ▶ When a process **gets all its resources**, it must **return them** in a **finite amount of time**.



Data Structures for Banker's Algorithm

- ▶ n = number of processes, and m = number of resources types



Data Structures for Banker's Algorithm

- ▶ n = number of processes, and m = number of resources types
- ▶ *Available*: vector of length m .
 - If $Available[j] = k$, there are k instances of resource type R_j available.

Data Structures for Banker's Algorithm

- ▶ n = number of processes, and m = number of resources types
- ▶ *Available*: vector of length m .
 - If $Available[j] = k$, there are k instances of resource type R_j available.
- ▶ *Max*: $n \times m$ matrix.
 - If $Max[i, j] = k$, then process P_i may request at most k instances of resource type R_j .

Data Structures for Banker's Algorithm

- ▶ n = number of processes, and m = number of resources types
- ▶ *Available*: vector of length m .
 - If $Available[j] = k$, there are k instances of resource type R_j available.
- ▶ *Max*: $n \times m$ matrix.
 - If $Max[i, j] = k$, then process P_i may request at most k instances of resource type R_j .
- ▶ *Allocation*: $n \times m$ matrix.
 - If $Allocation[i, j] = k$ then P_i is currently allocated k instances of R_j .

Data Structures for Banker's Algorithm

- ▶ n = number of processes, and m = number of resources types
- ▶ *Available*: vector of length m .
 - If $Available[j] = k$, there are k instances of resource type R_j available.
- ▶ *Max*: $n \times m$ matrix.
 - If $Max[i, j] = k$, then process P_i may request at most k instances of resource type R_j .
- ▶ *Allocation*: $n \times m$ matrix.
 - If $Allocation[i, j] = k$ then P_i is currently allocated k instances of R_j .
- ▶ *Need*: $n \times m$ matrix.
 - If $Need[i, j] = k$, then P_i may need k more instances of R_j to complete its task $Need[i, j] = Max[i, j] - Allocation[i, j]$



Safety Algorithm

1. Let *Work* and *Finish* be vectors of length m and n , respectively.

Initialize:

Work = *Available*

Finish[i] = *false* for $i = 0, 1, \dots, n - 1$



Safety Algorithm

1. Let *Work* and *Finish* be vectors of length m and n , respectively.

Initialize:

Work = *Available*

Finish[i] = *false* for $i = 0, 1, \dots, n - 1$

2. Find an i such that both:
 1. *Finish*[i] = *false*
 2. $Need_i \leq Work$If no such i exists, go to step 4.

Safety Algorithm

1. Let *Work* and *Finish* be vectors of length m and n , respectively.

Initialize:

Work = *Available*

Finish[i] = *false* for $i = 0, 1, \dots, n - 1$

2. Find an i such that both:

1. *Finish*[i] = *false*

2. $Need_i \leq Work$

If no such i exists, go to step 4.

3. *Work* = *Work* + *Allocation* _{i}

Finish[i] = *true*

Go to step 2

Safety Algorithm

1. Let *Work* and *Finish* be vectors of length m and n , respectively.

Initialize:

Work = *Available*

Finish[i] = *false* for $i = 0, 1, \dots, n - 1$

2. Find an i such that both:

1. *Finish*[i] = *false*

2. $Need_i \leq Work$

If no such i exists, go to step 4.

3. *Work* = *Work* + *Allocation* _{i}

Finish[i] = *true*

Go to step 2

4. If *Finish*[i] == *true* for all i , then the system is in a **safe state**.



Resource-Request Algorithm for Process P_i (1/2)

- ▶ *Request_i* = request vector for process P_i . If $Request_i[j] = k$, then process P_i wants k instances of resource type R_j .

Resource-Request Algorithm for Process P_i (1/2)

- ▶ $Request_i$ = request vector for process P_i . If $Request_i[j] = k$, then process P_i wants k instances of resource type R_j .
- ▶ 1. If $Request_i \leq Need_i$, go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.

Resource-Request Algorithm for Process P_i (1/2)

- ▶ $Request_i$ = request vector for process P_i . If $Request_i[j] = k$, then process P_i wants k instances of resource type R_j .
- ▶ 1. If $Request_i \leq Need_i$, go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
- ▶ 2. If $Request_i \leq Available$, go to step 3. Otherwise P_i must wait, since resources are not available.

Resource-Request Algorithm for Process P_i (2/2)

- ▶ 3. Pretend to allocate requested resources to P_i by modifying the state as follows:

$$Available = Available - Request_i$$

$$Allocation_i = Allocation_i + Request_i$$

$$Need_i = Need_i - Request_i$$

Resource-Request Algorithm for Process P_i (2/2)

- ▶ 3. Pretend to allocate requested resources to P_i by modifying the state as follows:

$Available = Available - Request_i$

$Allocation_i = Allocation_i + Request_i$

$Need_i = Need_i - Request_i$

- If **safe**: the resources are allocated to P_i

Resource-Request Algorithm for Process P_i (2/2)

- 3. Pretend to allocate requested resources to P_i by modifying the state as follows:

$Available = Available - Request_i$

$Allocation_i = Allocation_i + Request_i$

$Need_i = Need_i - Request_i$

- If **safe**: the resources are allocated to P_i
- If **unsafe**: P_i must wait, and the old resource-allocation state is restored

Banker's Algorithm Example (1/3)

- ▶ 5 processes: P_0 through P_4
- ▶ 3 resource types:
 - A (10 instances), B (5 instances), and C (7 instances)
- ▶ Snapshot at time T_0

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
P_0	0 1 0	7 5 3	3 3 2
P_1	2 0 0	3 2 2	
P_2	3 0 2	9 0 2	
P_3	2 1 1	2 2 2	
P_4	0 0 2	4 3 3	

Banker's Algorithm Example (2/3)

- The content of the matrix *Need* is defined to be $Max - Allocation$

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>		<u>Need</u>
	A B C	A B C	A B C		A B C
P_0	0 1 0	7 5 3	3 3 2	P_0	7 4 3
P_1	2 0 0	3 2 2		P_1	1 2 2
P_2	3 0 2	9 0 2		P_2	6 0 0
P_3	2 1 1	2 2 2		P_3	0 1 1
P_4	0 0 2	4 3 3		P_4	4 3 1

Banker's Algorithm Example (2/3)

- The content of the matrix *Need* is defined to be $Max - Allocation$

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>		<u>Need</u>
	A B C	A B C	A B C		A B C
P_0	0 1 0	7 5 3	3 3 2	P_0	7 4 3
P_1	2 0 0	3 2 2		P_1	1 2 2
P_2	3 0 2	9 0 2		P_2	6 0 0
P_3	2 1 1	2 2 2		P_3	0 1 1
P_4	0 0 2	4 3 3		P_4	4 3 1

- Is the system safe?

Banker's Algorithm Example (2/3)

- The content of the matrix *Need* is defined to be $Max - Allocation$

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>		<u>Need</u>
	A B C	A B C	A B C		A B C
P_0	0 1 0	7 5 3	3 3 2	P_0	7 4 3
P_1	2 0 0	3 2 2		P_1	1 2 2
P_2	3 0 2	9 0 2		P_2	6 0 0
P_3	2 1 1	2 2 2		P_3	0 1 1
P_4	0 0 2	4 3 3		P_4	4 3 1

- Is the system safe? $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ satisfies safety criteria.



Banker's Algorithm Example (3/3)

- ▶ P_1 Request $(1, 0, 2)$
- ▶ Check that $Request \leq Available$: $(1, 0, 2) \leq (3, 3, 2) \Rightarrow true$

Banker's Algorithm Example (3/3)

- ▶ P_1 Request $(1, 0, 2)$
- ▶ Check that $Request \leq Available$: $(1, 0, 2) \leq (3, 3, 2) \Rightarrow true$

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	7 4 3	2 3 0
P_1	3 0 2	0 2 0	
P_2	3 0 2	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

Banker's Algorithm Example (3/3)

- ▶ P_1 Request $(1, 0, 2)$
- ▶ Check that $Request \leq Available$: $(1, 0, 2) \leq (3, 3, 2) \Rightarrow true$

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	7 4 3	2 3 0
P_1	3 0 2	0 2 0	
P_2	3 0 2	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

- ▶ Executing **safety algorithm** shows that sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies safety requirement.

Banker's Algorithm Example (3/3)

- ▶ P_1 Request $(1, 0, 2)$
- ▶ Check that $Request \leq Available$: $(1, 0, 2) \leq (3, 3, 2) \Rightarrow true$

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	7 4 3	2 3 0
P_1	3 0 2	0 2 0	
P_2	3 0 2	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

- ▶ Executing **safety algorithm** shows that sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies safety requirement.
- ▶ Can request for $(3, 3, 0)$ by P_4 be granted?

Banker's Algorithm Example (3/3)

- ▶ P_1 Request $(1, 0, 2)$
- ▶ Check that $Request \leq Available$: $(1, 0, 2) \leq (3, 3, 2) \Rightarrow true$

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	7 4 3	2 3 0
P_1	3 0 2	0 2 0	
P_2	3 0 2	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

- ▶ Executing **safety algorithm** shows that sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies safety requirement.
- ▶ Can request for $(3, 3, 0)$ by P_4 be granted?
- ▶ Can request for $(0, 2, 0)$ by P_0 be granted?

Deadlock Detection



Deadlock Detection

- ▶ Allow system to enter deadlock state
- ▶ Detection algorithm
- ▶ Recovery scheme



Single Instance of Each Resource Type

- ▶ Maintain **wait-for** graph.
 - **Nodes** are **processes**.
 - $P_i \rightarrow P_j$ if P_i is **waiting** for P_j .

Single Instance of Each Resource Type

- ▶ Maintain **wait-for** graph.
 - **Nodes** are **processes**.
 - $P_i \rightarrow P_j$ if P_i is **waiting** for P_j .

- ▶ Periodically invoke an algorithm that searches for a **cycle in the graph**.


Single Instance of Each Resource Type

- ▶ Maintain **wait-for** graph.
 - **Nodes** are **processes**.
 - $P_i \rightarrow P_j$ if P_i is **waiting** for P_j .
- ▶ Periodically invoke an algorithm that searches for a **cycle in the graph**.
- ▶ If there is a **cycle**, there exists a **deadlock**.


Single Instance of Each Resource Type

- ▶ Maintain **wait-for** graph.
 - **Nodes** are **processes**.
 - $P_i \rightarrow P_j$ if P_i is **waiting** for P_j .
- ▶ Periodically invoke an algorithm that searches for a **cycle in the graph**.
- ▶ If there is a **cycle**, there exists a **deadlock**.
- ▶ An algorithm to **detect a cycle in a graph** requires an $O(n^2)$ operations, where n is the number of **vertices** in the graph.

Resource-Allocation Graph and Wait-for Graph



Resource-allocation graph



Corresponding Wait-for graph



Data Structures for Deadlock Detection

- ▶ *Available*: vector of length m , indicates the number of **available resources** of each type.



Data Structures for Deadlock Detection

- ▶ *Available*: vector of length m , indicates the number of **available resources** of each type.
- ▶ *Allocation*: $n \times m$ matrix, defines the **number of resources** of each type currently **allocated** to each process.

Data Structures for Deadlock Detection

- ▶ *Available*: vector of length m , indicates the number of **available resources** of each type.
- ▶ *Allocation*: $n \times m$ matrix, defines the **number of resources** of each type currently **allocated** to each process.
- ▶ *Request*: $n \times m$ matrix, indicates the **current request** of each process.
 - If $Request[i, j] = k$, then P_i **requesting** k more instances of resource type R_j .



Detection Algorithm (1/2)

- ▶ 1. Let *Work* and *Finish* be vectors of length m and n , respectively.
Initialize:
 - a. $Work = Available$
 - b. For $i = 1, 2, \dots, n$, if $Allocation_i \neq 0$, then $Finish[i] = false$;
otherwise, $Finish[i] = true$

Detection Algorithm (1/2)

- ▶ 1. Let *Work* and *Finish* be vectors of length m and n , respectively.
Initialize:
 - a. $Work = Available$
 - b. For $i = 1, 2, \dots, n$, if $Allocation_i \neq 0$, then $Finish[i] = false$;
otherwise, $Finish[i] = true$

- ▶ 2. Find an index i such that both:
 - a. $Finish[i] == false$
 - b. $Request_i \leq Work$

Detection Algorithm (1/2)

- ▶ 1. Let *Work* and *Finish* be vectors of length m and n , respectively.
Initialize:
 - a. $Work = Available$
 - b. For $i = 1, 2, \dots, n$, if $Allocation_i \neq 0$, then $Finish[i] = false$;
otherwise, $Finish[i] = true$

- ▶ 2. Find an index i such that both:
 - a. $Finish[i] == false$
 - b. $Request_i \leq Work$

- ▶ If no such i exists, go to step 4



Detection Algorithm (2/2)

- ▶ 3. $Work = Work + Allocation;$
 $Finish[i] = true$
go to step 2

Detection Algorithm (2/2)

- ▶ 3. $Work = Work + Allocation_i$
 $Finish[i] = true$
go to step 2
- ▶ 4. If $Finish[i] == false$, for some i , $1 \leq i \leq n$, then the system is in **deadlock** state. Moreover, if $Finish[i] == false$, then P_i is **deadlocked**.

Detection Algorithm (2/2)

- ▶ 3. $Work = Work + Allocation_i$
 $Finish[i] = true$
go to step 2
- ▶ 4. If $Finish[i] == false$, for some i , $1 \leq i \leq n$, then the system is in **deadlock** state. Moreover, if $Finish[i] == false$, then P_i is **deadlocked**.
- ▶ Algorithm requires an order of $O(m \times n^2)$ operations to detect whether the system is in deadlocked state.

Detection Algorithm Example (1/2)

- ▶ 5 processes: P_0 through P_4
- ▶ 3 resource types:
 - A (7 instances), B (2 instances), and C (6 instances)
- ▶ Snapshot at time T_0

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
P_0	0 1 0	0 0 0	0 0 0
P_1	2 0 0	2 0 2	
P_2	3 0 3	0 0 0	
P_3	2 1 1	1 0 0	
P_4	0 0 2	0 0 2	

Detection Algorithm Example (1/2)

- ▶ 5 processes: P_0 through P_4
- ▶ 3 resource types:
 - A (7 instances), B (2 instances), and C (6 instances)
- ▶ Snapshot at time T_0

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
P_0	0 1 0	0 0 0	0 0 0
P_1	2 0 0	2 0 2	
P_2	3 0 3	0 0 0	
P_3	2 1 1	1 0 0	
P_4	0 0 2	0 0 2	

- ▶ Deadlock?

Detection Algorithm Example (1/2)

- ▶ 5 processes: P_0 through P_4
- ▶ 3 resource types:
 - A (7 instances), B (2 instances), and C (6 instances)
- ▶ Snapshot at time T_0

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
P_0	0 1 0	0 0 0	0 0 0
P_1	2 0 0	2 0 2	
P_2	3 0 3	0 0 0	
P_3	2 1 1	1 0 0	
P_4	0 0 2	0 0 2	

- ▶ **Deadlock?** Sequence $\langle P_0, P_2, P_3, P_1, P_4 \rangle$ will result in $Finish[i] = true$ for all i

Detection Algorithm Example (2/2)

- P_2 requests an additional instance of type C

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>		<u>Request</u>
	$A\ B\ C$	$A\ B\ C$	$A\ B\ C$		$A\ B\ C$
P_0	0 1 0	0 0 0	0 0 0	P_0	0 0 0
P_1	2 0 0	2 0 2		P_1	2 0 2
P_2	3 0 3	0 0 0		P_2	0 0 1
P_3	2 1 1	1 0 0		P_3	1 0 0
P_4	0 0 2	0 0 2		P_4	0 0 2

Detection Algorithm Example (2/2)

- ▶ P_2 requests an additional instance of type C

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>		<u>Request</u>
	$A\ B\ C$	$A\ B\ C$	$A\ B\ C$		$A\ B\ C$
P_0	0 1 0	0 0 0	0 0 0	P_0	0 0 0
P_1	2 0 0	2 0 2		P_1	2 0 2
P_2	3 0 3	0 0 0		P_2	0 0 1
P_3	2 1 1	1 0 0		P_3	1 0 0
P_4	0 0 2	0 0 2		P_4	0 0 2

- ▶ Can reclaim resources held by process P_0 , but **insufficient resources** to fulfill other processes; requests

Detection Algorithm Example (2/2)

- ▶ P_2 requests an additional instance of type C

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>		<u>Request</u>
	$A\ B\ C$	$A\ B\ C$	$A\ B\ C$		$A\ B\ C$
P_0	0 1 0	0 0 0	0 0 0	P_0	0 0 0
P_1	2 0 0	2 0 2		P_1	2 0 2
P_2	3 0 3	0 0 0		P_2	0 0 1
P_3	2 1 1	1 0 0		P_3	1 0 0
P_4	0 0 2	0 0 2		P_4	0 0 2

- ▶ Can reclaim resources held by process P_0 , but **insufficient resources** to fulfill other processes; requests
- ▶ **Deadlock** exists, consisting of processes P_1 , P_2 , P_3 , and P_4

Recovery From Deadlock



Recovery from Deadlock

- ▶ Process **termination**
- ▶ Resource **preemption**



Process Termination

- ▶ Abort **all deadlocked** processes.



Process Termination

- ▶ Abort **all deadlocked** processes.
- ▶ Abort **one process at a time** until the deadlock cycle is eliminated



Process Termination

- ▶ Abort **all deadlocked** processes.
- ▶ Abort **one process at a time** until the deadlock cycle is eliminated
- ▶ In which order should we choose to abort?

Process Termination

- ▶ Abort **all deadlocked** processes.
- ▶ Abort **one process at a time** until the deadlock cycle is eliminated
- ▶ In which order should we choose to abort?
 1. Priority of the process.
 2. How long process has computed, and how much longer to completion.
 3. Resources the process has used.
 4. Resources process needs to complete.
 5. How many processes will need to be terminated.
 6. Is process interactive or batch.



Resource Preemption

- ▶ Selecting a victim: minimize cost



Resource Preemption

- ▶ Selecting a **victim**: minimize cost
- ▶ **Rollback**: return to some **safe state**, restart process for that state.



Resource Preemption

- ▶ Selecting a **victim**: minimize cost
- ▶ **Rollback**: return to some **safe state**, restart process for that state.
- ▶ **Starvation**: same process may always be picked as victim, include number of rollback in cost factor.

Summary



Summary

- ▶ Deadlock



Summary

- ▶ Deadlock
- ▶ Four simultaneous conditions: mutual exclusion, hold and wait, no pre-emption, circular wait



Summary

- ▶ Deadlock
- ▶ Four simultaneous conditions: mutual exclusion, hold and wait, no pre-emption, circular wait
- ▶ Deadlock prevention



Summary

- ▶ Deadlock
- ▶ Four simultaneous conditions: mutual exclusion, hold and wait, no pre-emption, circular wait
- ▶ Deadlock prevention
- ▶ Deadlock avoidance: resource-allocation algorithm, banker's algorithm

Summary

- ▶ Deadlock
- ▶ Four simultaneous conditions: mutual exclusion, hold and wait, no pre-emption, circular wait
- ▶ Deadlock prevention
- ▶ Deadlock avoidance: resource-allocation algorithm, banker's algorithm
- ▶ Deadlock detection: Wait-for graph

Summary

- ▶ Deadlock
- ▶ Four simultaneous conditions: mutual exclusion, hold and wait, no pre-emption, circular wait
- ▶ Deadlock prevention
- ▶ Deadlock avoidance: resource-allocation algorithm, banker's algorithm
- ▶ Deadlock detection: Wait-for graph
- ▶ Deadlock recovery: process termination, resource preemption

Questions?

Acknowledgements

Some slides were derived from Avi Silberschatz slides.